

Introduction of Reinforcement Learning

Tian Lan

2018-11-3

- ① David Silver UCL course
- ② My notebook for RL
- ③ One standard repo for RL agent in PyTorch
- ④ Alibaba Reinforcement Learning Application
- ⑤ Lots of repos can be found on GitHub

- ① Introduction and MDP
- ② Planning by DP
- ③ Model-Free Value-Based
- ④ Value function approximation
 - ① FA and its limitation
 - ② DQN and extensions
- ⑤ Model-Free Policy-Based
 - ① Policy Gradient
 - ② Algorithms
- ⑥ Model-Based
 - ① Integrating Learning
 - ② Dyna-Q
- ⑦ Application
 - ① RL in NLP
 - ② Alibaba StarCraft

1. Introduction




Figure 1: Reinforcement Learning

1. Introduction and MDP

Reinforcement Learning:

Dynamic decision process with delay reward

- ① no supervisor, only reward signal from environment
- ② Feedback is delayed
- ③ Time really matters
- ④ Agent's action affect subsequent data it receives

Reward hypothesis:

All goals can be described by Maximisation of expected cumulative rewards

RL Task

Prediction and Control

1. Introduction and MDP

What is MDP and why we need it?

- ① Markov decision processes formally describe the Environment for RL
In RL, we only care about Agent and Environment
- ② Current state is known, the history can be thrown away

$$S_t = F(h_{<t})$$

- ③ Almost all RL problems can be formalised as MDPs

1. Introduction and MDP

① MDP definition

$$(S, A, R, P, \gamma)$$

$$P_{ss'}^a = \mathbb{P}[S_{t+1} = s' | S_t = s]$$

$$R_s^a = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

$$\gamma \in [0, 1]$$

② Policy definition

Define the distribution of the action condition on state

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s]$$

③ Return Reward

Cumulative reward from current state

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+1+k}$$

1. Introduction and MDP

① Value function

- State Value function

$$v_{\pi}(s) = \mathbb{E}[G_t | S_t = s]$$

- Action Value function

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a]$$

② Bellman Equation

- State Value Bellman function

$$v_{\pi}(s) = \mathbb{E}_{\pi}[R_{t+1} + \gamma v_{\pi}(S_{t+1}) | S_t = s]$$

- Action Value Bellman function

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1}) | S_t = s, A_t = a]$$

1. Introduction and MDP

Optimal Bellman Equation

$$v_*(s) = \max_a q_*(s, a)$$

$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s')$$

$$q_*(s, a) = R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a \max_{a'} q_*(s', a')$$

Find stable $v_*(s)$, prediction problem solved

Find stable $p_*(s, a)$, control problem solved

Need to use iterate the equation, because the Optimal Bellman Equation is nonlinear.

2. Planning by DP

Dynamic Programming: Optimal method for decision process

- Optimal substructure
Optimal solution can be decomposed into subproblems
- Overlapping subproblems
Solutions can be cached and reused
- Markov Decision Process satisfy the properties
 - Bellman equation gives the way to decompose solution
 - Value function store solution

2. Planning by DP

Prediction

- Problem: MDP and policy π given, solve v_π , evaluate policy π
- Solution: iterative application of Bellman equation

$$v_1 \rightarrow v_2 \rightarrow \dots \rightarrow v_\pi$$

- Convergence of v_π can be proven

$$v_{k+1}(s) = \sum_{a \in A} \pi(a|s) (R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_k(s'))$$

$$v^{k+1} = R^\pi + \gamma P^\pi v^k$$

2. Planning by DP

Control

- Problem: MDP and initial policy π given, return best π
- Solution: Policy Iteration


- Evaluation

$$v_{\pi}(s) = \mathbb{E}[R_{t+1} + \gamma R_{t+2} + \dots | S_t = s]$$

- Improve policy

$$\pi' = greedy(v_{\pi})$$

2. Planning by DP



Policy evaluation Estimate v_π

Iterative policy evaluation

Policy improvement Generate $\pi' \geq \pi$

Greedy policy improvement




Figure 2: Policy Iteration

2. Planning by DP

- Convergence

$$v_{\pi}(s) = \max_{a \in A} q_{\pi}(s, a) = v_{*}(s)$$

$$q_{*}(s, a) = R_s^a + \gamma \sum_{s' \in S'} P_{ss'}^a v_{*}(s')$$

$$v_{\pi}(s) = q_{\pi}(s, \pi'(s)) = \max_{a \in A} q_{\pi}(s, a) \geq q_{\pi}(s, \pi(s)) = v_{\pi}(s)$$

- Optimal Bellman equation satisfy, solve the MDP control problem

2. Planning by DP

- Value Iteration for control problem
Iterative application of this equation

$$v_*(s) \leftarrow \max_{a \in A} R_s^a + \gamma \sum_{s' \in S} P_{ss'}^a v_*(s')$$

Value Iteration is another way to solve the control and prediction problem

- Constraint for Planning by DP
MDP must be given, which means that $P_{ss'}^a$ and R_s^a can be accessed.

3. Model-Free Value-Based

Definition

- ① MDP is unknown, but we still want to solve it. $P_{ss'}^a$ and R_s^a is unknown to us.
- ② We can only learn the value or policy from the experience.
- ③ Find the optimal V value function, prediction solved.
- ④ Find the optimal Q value function, control solved.

3. Model-Free Prediction

Monte-Carlo Learning

- ① Learn from the episode experience with environment
- ② MC is model-free: no knowledge of MDP (P, R)
- ③ MC learns from complete episode without bootstrapping
- ④ MC's idea: $value = \mathbb{E}[G]$
- ⑤ Constraint: All episodes must terminate

3. Model-Free Prediction

Monte-Carlo Learning

- ① Goal: learn v_π from episodes of experience under policy π

$$S_1, A_1, R_2, \dots, S_k \sim \pi$$

- ② Discounted reward of episode

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

- ③ Update $v(s)$ after episode

$$N(S_t) \leftarrow N(S_t) + 1$$

$$V(S_t) \leftarrow V(S_t) + \frac{1}{N(S_t)} (G_t - V(S_t))$$

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$

3. Model-Free Prediction

Temporal-Difference Learning

- ① TD methods learn directly from episodes of experience
- ② TD is model-free: no knowledge of MDP (P, R)
- ③ TD learns from incomplete episode with bootstrapping
- ④ TD update the value function with the guess from previous iteration

3. Model-Free Prediction

Temporal-Difference Learning

- ① Goal: learn v_π online from experience under policy π
- ② MC update

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t))$$

- ③ TD(0)

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$

- $R_{t+1} + \gamma V(S_{t+1})$ is called TD target
- $\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$ is called TD error

3. Model-Free Prediction

Comparison

- 1 TD learns online but MC wait until end of episode
- 2 MC has high variance and zero bias, TD has low variance and some bias
- 3 Convergence:
 - 1 MC converges to solution with MSE

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(s_t^k))^2$$

- 2 TD converges to max likelihood Markov model
TD is better suited to the Markov environment.

3. Model-Free Prediction


TD(λ)

$$E_0(s) = 0$$

$$E_t(s) = \gamma\lambda E_{t-1}(s) + 1(\text{if } S_t = s)$$

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$

$$V(s) \leftarrow V(s) + \alpha \delta_t E_t(s)$$



- TD(λ) converges faster than TD(0)
- δ_t define the result of fixing
- E_t define the degree of fixing

3. Model-Free Control

Generalised Policy Iteration




Figure 3: Generalised Policy Iteration

3. Model-Free Control

Sarsa

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
Repeat (for each episode):  
  Initialize  $S$   
  Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
  Repeat (for each step of episode):  
    Take action  $A$ , observe  $R, S'$   
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
     $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$   
     $S \leftarrow S'; A \leftarrow A';$   
  until  $S$  is terminal
```

Figure 4: Sarsa Algorithm

3. Model-Free Control

Sarsa(λ)

```
Initialize  $Q(s, a)$  arbitrarily, for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
Repeat (for each episode):
   $E(s, a) = 0$ , for all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ 
  Initialize  $S, A$ 
  Repeat (for each step of episode):
    Take action  $A$ , observe  $R, S'$ 
    Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)
     $\delta \leftarrow R + \gamma Q(S', A') - Q(S, A)$ 
     $E(S, A) \leftarrow E(S, A) + \delta$ 
    For all  $s \in \mathcal{S}, a \in \mathcal{A}(s)$ :
       $Q(s, a) \leftarrow Q(s, a) + \alpha \delta E(s, a)$ 
       $E(s, a) \leftarrow \gamma \lambda E(s, a)$ 
     $S \leftarrow S'; A \leftarrow A'$ 
  until  $S$  is terminal
```

Figure 5: Sarsa(λ) Algorithm

3. Model-Free Control

Off-Policy and On-Policy

① On-Policy

- learn about policy π from experience sampled from π
- The special case of the Off-Policy

② Off-Policy

- learn about policy π from experience sampled from μ
- Re-use experience generated from other policy (other agent or even human)
- Balance between Exploitation and Exploration through Q-Learning

3. Model-Free Control

Q-Learning

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
Repeat (for each episode):  
  Initialize  $S$   
  Repeat (for each step of episode):  
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
    Take action  $A$ , observe  $R, S'$   
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$   
     $S \leftarrow S'$ ;  
  until  $S$  is terminal
```

Figure 6: Off-Policy TD Learning, Q-Learning

4. Value Function Approximation

- ① Reinforcement Learning can be used to solve large problem
 - Backgammon: 10^{20} states
 - Game of Go: 10^{170} states
 - Some environment: continuous state space
- ② Problem:
 - Tabular method need lots of memory to store Value Function
 - Learn value individually slowly
- ③ Solution - **Value Function Approximation**
 - Save the memory with approximation parameters w

$$v(s, w) \approx v_{\pi}(s)$$

$$q(s, a, w) \approx q_{\pi}(s, a)$$

- Extension to unseen states

4. Value Function Approximation

① Approximate model

- Nonlinear model: Nerual Networks
- Linear regression
- Other regression model ...

② Optimal parameters

- Gradient descent
- Batch training

4. Value Function Approximation

Gradient Descent

- ① Goal: Find the parameter vector w minimising MSE between approximate value function $v(s, w)$ and true value function $v_\pi(s)$

$$J(w) = \mathbb{E}_\pi[(v_\pi(s) - v(s, w))^2]$$

- ② Gradient descent

$$\delta w = -\frac{1}{2}\alpha \nabla_w J(w) = \alpha \mathbb{E}_\pi[(v_\pi(s) - v(s, w)) \nabla_w v(s, w)]$$

$$\delta w = \alpha (v_\pi(s) - v(s, w)) \nabla_w v(s, w)$$

- ③ SGD \rightarrow Expectation

4. Value Function Approximation

Gradient Descent

- ① Linear function, $x(S)$ is the feature vector

$$v(s, w) = x(s)^T w$$

- ② Object function

$$J(w) = \mathbb{E}_{\pi}[(v_{\pi}(s) - x(s)^T w)^2]$$

- ③ Gradient Descent

$$\delta w = \alpha(v_{\pi}(s) - v(s, w))x(s)$$

- ④ Lookup table is the special case of Linear function

$$v(s, w) = x_{one-hot}^T w$$

4. Value Function Approximation

Gradient Descent

① MC

$$\delta w = \alpha(G_t - v(s_t, w))\nabla_w v(s_t, w)$$

② TD(0)

$$\delta w = \alpha(R_{t+1} + \gamma v(s_{t+1}, w) - v(s_t, w))\nabla_w v(s_t, w)$$

③ TD(λ)

$$\delta w = \alpha(G_t^\lambda - v(s, w))\nabla_w v(s_t, w)$$

$$\delta_t = R_{t+1} + \gamma v(s_{t+1}, w) - v(s_t, w)$$

$$E_t = \gamma\lambda E_{t-1} + \delta_t$$

$$\delta w = \alpha\delta_t E_t$$

4. Value Function Approximation

Limitation

Algorithm	Table Lookup	Linear	Non-Linear
Monte-Carlo Control	✓	(✓)	✗
Sarsa	✓	(✓)	✗
Q-learning	✓	✗	✗
Gradient Q-learning	✓	✓	✗

(✓) = chatters around near-optimal value function

Figure 7: Neural Networks don't work¹

An analysis of temporal-difference learning with function approximation.
IEEE 1997

4. Value Function Approximation

Batch Methods - DQN

① DeepMind Nature 2015

- DRL of DQN
- integrate the CNN and Q-Learning to master the games
- Solve the divergence of nonlinear function approximation
Q-Learning

② Idea

- Experience replay: avoid relevance
- Fixed Q target Network: Off-Policy

4. Value Function Approximation

Algorithm 1 Deep Q-learning with experience replay

Initialize replay memory D to capacity N
Initialize action-value function Q with random weights θ
Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$
for episode 1, M **do** Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$
 for $t = 1, T$ **do**
 With probability ε select a random action a_t
 otherwise select $a_t = \arg \max_a Q(\phi(s_t), a; \theta)$
 Execute action a_t in the emulator and observe reward r_t and image x_{t+1}
 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$
 Store experience $(\phi_t, a_t, r_t, \phi_{t+1})$ in D
 Sample random minibatch of experiences $(\phi_j, a_j, r_j, \phi_{j+1})$ from D
 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$
 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the weights θ
 Every C steps reset $\hat{Q} = Q$
 end for
end for

Figure 8: DQN

4. Value Function Approximation

Extension of DQN

① Double DQN

- fix the overestimation problem of Q-Learning
- avoid suboptimal result

$$loss = (r + \gamma Q(s', \arg \max_a Q(s', a', w), w^-) - Q(s, a, w))^2$$

② Prioritized Replay DQN

- When reward is sparse, DQN is hard to train
- TD error is the weight

③ Dueling DQN

- For some states, V function is more important
- speed up

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$$

5. Policy Gradient

Value-Based and Policy-Based RL

- Value Based
 - Learn value function
 - implicit policy (such as ϵ -greedy)
- Policy Based
 - No value function
 - Learn policy directly
- Actor-Critic
 - Critic: Learn Value Function
 - Actor: Learn policy with the guidance of Critic

5. Policy Gradient




Figure 9: Aliased

5. Policy Gradient

Advantage

- Learn policy with probability
- Learn for continuous state space, better than Value-Based

5. Policy Gradient

- 1 Parameterise the **policy**

$$\pi_{\theta}(s, a) = \mathbb{P}[a|s, \theta]$$

- 2 Evaluate the policy with object function

$$J(\theta) = \sum \pi_{\theta}(a|s) Q(s, a; \theta)$$

$$J(\theta) = \mathbb{E}_{\pi_{\theta}, s \sim \rho^{\pi}}[Q(s, a; \theta)]$$

- 3 Optimize

- 1 Gradient Ascent
- 2 No Gradient Ascent

5. Policy Gradient

Compute the score function for gradient

- 1 Discrete environment - **Softmax Policy**

$$\pi_{\theta}(s, a) \propto e^{\phi(s, a)^T \theta}$$

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \phi(s, a) - \mathbb{E}_{\theta}[\phi(s, \cdot)]$$

$$\nabla_{\theta} \pi_{\theta}(s, a) = \pi_{\theta}(s, a) \frac{\nabla_{\theta} \pi_{\theta}(s, a)}{\pi_{\theta}(s, a)} = \pi_{\theta}(s, a) \nabla_{\theta} \log \pi_{\theta}(s, a)$$

- 2 Continuing environment - **Gaussian Policy**

$$\nabla_{\theta} \log \pi_{\theta}(s, a) = \frac{(a - \mu(s))\phi(s)}{\sigma^2}$$

5. Policy Gradient

Policy Gradient Theorem

For any policy object function, the **Policy Gradient** is the same^a

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q^{\pi_{\theta}}(s, a)]$$

- ❶ $\nabla_{\theta} \log \pi_{\theta}(s, a)$ means the gradient
- ❷ Q means the degree of fixing
- ❸ Used to minimize the $-\log \pi_{\theta}(s, a) Q^{\pi_{\theta}}(s, a)$

^a<https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>

5. Policy Gradient Algorithms

MC Policy Gradient

```
function REINFORCE
  Initialise  $\theta$  arbitrarily
  for each spisode  $\{s_1, a_1, r_2, \dots\} \sim \pi_\theta$  do
    for  $t = 1$  to  $T - 1$  do
       $\theta \leftarrow \theta + \alpha \nabla_\theta \log \pi_\theta(s_t, a_t) v_t$ 
    end for
  end for
  return  $\theta$ 
end function
```

5. Policy Gradient Algorithms

TD Policy Gradient

```
function TD Policy Gradient
  Initialise  $\theta$  arbitrarily
  for each episode do
    for  $t = 1$  to  $T - 1$  do
       $\theta \leftarrow \theta + \alpha \nabla_{\theta} \log \pi_{\theta}(s_t, a_t) v_t$ 
    end for
  end for
  return  $\theta$ 
end function
```

5. Policy Gradient

Actor-Critic Policy Gradient

- ① MC has variance and need to be improved
- ② **Actor-Critic**
 - ① Critic: parametries $Q_w(s, a)$ for policy evaluation
 - ② Acotr: Policy Gradient with the guidance of Critic
 - ③ Result

$$\nabla_{\theta} J(\theta) \approx \mathbb{E}_{\pi_{\theta}} [\nabla_{\theta} \log \pi_{\theta}(s, a) Q_w(s, a)]$$

$$\delta \theta = \alpha \nabla_{\theta} \log \pi_{\theta}(s, a) Q_w(s, a)$$

5. Policy Gradient

Actor-Critic Policy Gradient with Linear function approximation

function QAC

 Initialise s, θ

 Sample $a \sim \pi_\theta$

 for each step do

 Sample reward $r = R_s^a$ and $s' \sim P_s^a$

 Sample action $a' \sim \pi_\theta(s', a')$

$\delta = r + \gamma Q_w(s', a') - Q_w(s, a)$

$\theta = \theta + \alpha \nabla_\theta \log \pi_\theta(s, a) Q_w(s, a)$

$w \leftarrow w + \beta \delta \phi(\cdot, sa)$

$a \leftarrow a', s \leftarrow s'$

 end for

end function

5. Policy Gradient

DDPG (Deep Deterministic Policy Gradient)

- DQN can't handle with high-dimension or continuous space
- Inspired by DQN, DeepMind propose DDPG
 - Actor-Critic: DQN is Critic and DPG² is Actor
 - Fixed Q target Network for policy and value
 - Target Network update slowly

²<http://proceedings.mlr.press/v32/silver14.pdf>

5. Policy Gradient

Algorithm 1 DDPG algorithm

Randomly initialize critic network $Q(s, a|\theta^Q)$ and actor $\mu(s|\theta^\mu)$ with weights θ^Q and θ^μ .

Initialize target network Q' and μ' with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\mu'} \leftarrow \theta^\mu$

Initialize replay buffer R

for episode = 1, M **do**

 Initialize a random process \mathcal{N} for action exploration

 Receive initial observation state s_1

for $t = 1, T$ **do**

 Select action $a_t = \mu(s_t|\theta^\mu) + \mathcal{N}_t$ according to the current policy and exploration noise

 Execute action a_t and observe reward r_t and observe new state s_{t+1}

 Store transition (s_t, a_t, r_t, s_{t+1}) in R

 Sample a random minibatch of N transitions (s_i, a_i, r_i, s_{i+1}) from R

 Set $y_i = r_i + \gamma Q'(s_{i+1}, \mu'(s_{i+1}|\theta^{\mu'}))|\theta^{Q'}$

 Update critic by minimizing the loss: $L = \frac{1}{N} \sum_i (y_i - Q(s_i, a_i|\theta^Q))^2$

 Update the actor policy using the sampled policy gradient:

$$\nabla_{\theta^\mu} J \approx \frac{1}{N} \sum_i \nabla_a Q(s, a|\theta^Q)|_{s=s_i, a=\mu(s_i)} \nabla_{\theta^\mu} \mu(s|\theta^\mu)|_{s_i}$$

 Update the target networks:

$$\theta^{Q'} \leftarrow \tau \theta^Q + (1 - \tau) \theta^{Q'}$$

$$\theta^{\mu'} \leftarrow \tau \theta^\mu + (1 - \tau) \theta^{\mu'}$$

end for
end for

Figure 10: DDPG

6. Model-Based

Comparison

- ① Model-Free
Learn value function and policy from the experience
- ② Model-Based
Learn MDP from the experience and get the policy and value function by planning.

6. Model-Based




Figure 11: Model-Based method

6. Model-Based

- ① Use any supervised Learning model to learn the MDP
- ② Anything we know now can be used to solve MDP
- ③ A direct way to solve the MDP but the bias is a disadvantage
- ④ Model is the agent's impression of environment

6. Model-Based

- 1 Parameter η represents the MDP

$$P_{\eta} \approx P$$

$$R_{\eta} \approx R$$

- 2 Goal: estimate model MDP from experience $\{s_1, a_1, r_2, \dots, s_T\}$
- 3 Supervised learning problem

$$S_t, A_t \rightarrow R_{t+1}, S_{t+1}$$

- 4 After learning the MDP, solved by
 - Planning (Policy Iteration, Value Iteration)
 - Model-Free (DQN, Policy Gradient)

6. Model-Based

Dyna

- ① Thinking: Use the model to generate infinite samples
This property is very useful for time-cost environment
- ② Simulation: Apply **model-free** RL to evaluation the samples
and fix the policy

6. Model-Based




Figure 12: Dyna

6. Model-Based

Integrated Architecture for Dyna-Q

```
Initialize  $Q(s, a)$  and  $Model(s, a)$  for all  $s \in \mathcal{S}$  and  $a \in \mathcal{A}(s)$ 
Do forever:
  (a)  $S \leftarrow$  current (nonterminal) state
  (b)  $A \leftarrow \epsilon$ -greedy( $S, Q$ )
  (c) Execute action  $A$ ; observe resultant reward,  $R$ , and state,  $S'$ 
  (d)  $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
  (e)  $Model(S, A) \leftarrow R, S'$  (assuming deterministic environment)
  (f) Repeat  $n$  times:
     $S \leftarrow$  random previously observed state
     $A \leftarrow$  random action previously taken in  $S$ 
     $R, S' \leftarrow Model(S, A)$ 
     $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma \max_a Q(S', a) - Q(S, A)]$ 
```

Figure 13: Dyna-Q, Q-Learning, DQN

7. Application

Conclusion

- ① Reward is the core for Reinforcement Learning
- ② If the problem can be thought as the MDP / POMDP, RL can be used
- ③ If the reward is clear, RL can be used
- ④ Sample data is important, no experience no RL
- ⑤ RL is difficult to train, be careful about parameters

7. Application

RL in NLP




Figure 14: Pipeline Task-oriented dialogue system

7. Application

RL in NLP

For task-oriented dialogue, reward is clear

- ① DQN for Policy Learning
- ② Deep Dyna-Q (Dyna-Q + DQN) for Policy Learning
- ③ REINFORCE for Policy Learning

7. Application




Figure 1: The illustration of SeqGAN. Left: D is trained over the real data and the generated data by G . Right: G is trained by policy gradient where the final reward signal is provided by D and is passed back to the intermediate action value via Monte Carlo search.

Figure 15: SeqGAN

7. Application

SeqGAN

- Model the sequence generation as MDP
- The first work: integrate GANs with RL to generate discrete elements
- Generator is LSTM Policy (RL Agent)
REINFORCE Algorithm is used for updating G
- Discriminator is CNN with MC Search

7. Application

Alibaba StarCraft



Figure 16: StarCraft

7. Application

Alibaba StarCraft

- ① Alibaba choose the StarCraft as the platform for RL research
 - Game is the wonderful platform for RL
 - Experience data is infinite
 - Reward is clear
- ② Difficulies
 - Real-Time strategy
 - Huge state and action space
 - Incomplete nformation games
 - Inference for spatial and sequential
 - Multi Agent

7. Application

Alibaba StarCraft

Why GAMES?

- ① Advertisement in Alimama
- ② Search in TaoBao
- ③ Intelligent agent for task-oriented dialogue (AliMeChat)
- ④ Alibaba RL sort strategy decision model for 11.11

End

Question and Answer